

CMPSC 473 Project 1

Generated by Doxygen 1.9.1

1 CMPSC 473 Project 1	1
1.1 Main Data Structures	1
1.1.1 Queues	1
1.1.1.1 arrival_queue	1
1.1.1.2 ready_queue	1
1.1.1.3 io_queue	1
1.1.1.4 semaphore's blocked_queues	1
1.1.2 ProgramControlBlock	2
1.1.3 Semaphore	2
1.2 Theory of Operation	2
1.2.1 The start_check() Function	2
1.2.2 The schedule() Function	2
1.2.3 The cpu_me() function	3
1.2.4 The io_me() function	3
1.2.5 The P() function	3
1.2.6 The V() function	3
1.2.7 The end_me() function	3
2 Data Structure Index	5
2.1 Data Structures	5
3 File Index	7
3.1 File List	7
4 Data Structure Documentation	9
4.1 PriorityQueueEntryS Struct Reference	9
4.1.1 Detailed Description	10
4.2 PriorityQueueS Struct Reference	10
4.2.1 Detailed Description	10
4.3 ProgramControlBlockS Struct Reference	11
4.3.1 Detailed Description	11
4.4 SemaphoreS Struct Reference	11
4.4.1 Detailed Description	12
4.5 thread_struct Struct Reference	12
5 File Documentation	13
5.1 interface.h File Reference	13
5.1.1 Detailed Description	14
5.2 scheduler.h File Reference	14
5.2.1 Detailed Description	16
5.2.2 Macro Definition Documentation	16
5.2.2.1 MLFQ_LEVELS	16
5.2.3 Function Documentation	17
5.2.3.1 add_to_ready_queue()	17

5.2.3.2 <code>clean_up()</code>	17
5.2.3.3 <code>next_arrival()</code>	18
5.2.3.4 <code>next_io_queue_completion()</code>	18
5.2.3.5 <code>pick_from_arrival_queue()</code>	19
5.2.3.6 <code>pick_from_io_queue()</code>	20
5.2.3.7 <code>pick_from_ready_queue()</code>	21
5.2.3.8 <code>PriorityQueue_dequeue()</code>	21
5.2.3.9 <code>PriorityQueue_destroy()</code>	22
5.2.3.10 <code>PriorityQueue_enqueue()</code>	23
5.2.3.11 <code>PriorityQueue_init()</code>	24
5.2.3.12 <code>PriorityQueue_is_empty()</code>	24
5.2.3.13 <code>PriorityQueue_peek()</code>	25
5.2.3.14 <code>request_cpu()</code>	26
5.2.3.15 <code>request_io()</code>	26
5.2.3.16 <code>request_P()</code>	27
5.2.3.17 <code>request_V()</code>	28
5.2.3.18 <code>schedule()</code>	28
5.2.3.19 <code>start_check()</code>	29
5.2.3.20 <code>tick()</code>	30
5.2.4 Variable Documentation	31
5.2.4.1 <code>arrival_queue</code>	31
5.2.4.2 <code>current_cpu_program</code>	31
5.2.4.3 <code>execution_lock</code>	31
5.2.4.4 <code>io_queue</code>	32
5.2.4.5 <code>ready_queue</code>	32
5.2.4.6 <code>scheduler_type</code>	32
5.2.4.7 <code>wall_clock</code>	32
Index	33

Chapter 1

CMPSC 473 Project 1

Jinyu Liu(liu.jinyu@psu.edu), Hongyu Guo(hbg5147@psu.edu)

1.1 Main Data Structures

1.1.1 Queues

The priority queue is the data structure all queues are based on. It is implemented on a linked list. It takes a priority argument in its enqueue function to allow flexibility. Entries with smaller priority numbers have higher priorities and ties are broken with the thread_id where lower ids have higher priority.

1.1.1.1 arrival_queue

The arrival_queue holds threads that will arrive in the future. Its entries are ordered by arrival time

1.1.1.2 ready_queue

The ready queue(s) holds threads that are waiting for the CPU. Its behavior changes depending on the type of CPU scheduler specified.

1.1.1.3 io_queue

The IO queue holds the threads that are performing or waiting to perform IO. It is FCFS and the head of the queue is the thread currently doing IO.

1.1.1.4 semaphore's blocked_queues

The blocked queue holds threads that are blocked by this semaphore. It is FCFS.

1.1.2 ProgramControlBlock

The ProgramControlBlock are the block that contain information about each thread in the system. Most importantly, the `allow_to_run` conditional that is used to control the thread's execution.

1.1.3 Semaphore

Contain the value of semaphore and queue of threads blocked by the semaphore.

1.2 Theory of Operation

The program mostly runs single-threaded with only one thread actively running at any time to make sure to race conditions occurs while operating on the various data structures and to make scheduling easy. Each thread is blocked using its own conditional in its ProgramControlBlock, this allows the schedule function to choose which thread to execute next by signaling only the conditional blocking the specific thread. Once one thread is unblocked it runs until another call to schedule where control is transferred to another thread.

1.2.1 The start_check() Function

It is the job of the `start_check()` function to get the disorganized threads into this state where only one is running at a time. It enqueues the calling thread onto the arrival queue and blocks the calling thread on its conditional. The once the last thread calls `start_check`, `start_check` will make the first call to schedule and set in motion the scheduling operations.

1.2.2 The schedule() Function

The schedule function is heart of the scheduler. It decides which thread to transfer control to. The main goal of the schedule function is ensure that the `wall_clock` is not advanced until every thread has completed all operations in the current tick. To do this it priorities operations that do not advance the wall clock. It choose threads in the following order of priority.

1. Any thread in the `arrival_queue` that arrived at the current time.
2. Any thread in the `io_queue` that completed IO at the current time.
3. The thread in the `ready_queue` that is next in line for the CPU.

Once the schedule function finds the next thread to run according to the priority order above, It will transfer control to that thread.

However it is possible that the CPU is idle and all threads are either blocked or will arrive in the future. In this case we would need to skip time forward. Since threads blocked by semaphores require other threads to execute to be come unstuck, we don't need to consider them when figuring out how much to skip time forward by. Thus we only need to consider the following.

- The next thread in the `arrival_queue` that will arrive at some point in the future
- The thread currently doing IO that will complete at some point in the future

If the schedule function can not find any thread to run at the current time, it will find the earlier of the two from the list above, skip time to when that event happens, and transfer control to that thread. At this point schedule should have found a thread to give control to, if it still can't find a thread, it means there is a deadlock or bug, so it will raise an assertion error and terminate the program.

1.2.3 The `cpu_me()` function

It puts the thread on the `ready_queue` and then calls `schedule`. Once the scheduler decides it should have the CPU it will increment the wall clock by 1 by calling `tick(1)` to represent the 1 unit of time passes while it is using the CPU. it will return the current `wall_time` as required.

1.2.4 The `io_me()` function

It puts the thread on the `io_queue` and then calls `schedule`. Once the IO is completed, the scheduler will release the thread, and it will return the current `wall_clock`.

1.2.5 The `P()` function

It will decrement the value of the specified semaphore. If the value is negative, it will put the calling thread on the `blocked_queue` of the semaphore and call `schedule` to transfer control to another thread. After either it deciding not to block or the thread gets released by a corresponding call to `V()`, it will return the current `wall_time`.

1.2.6 The `V()` function

It will increment the value of the specified semaphore. If there are thread being blocked, it will moved one thread from `blocked_queue` to `arrival_queue`, to be run next time `schedule` is called.

1.2.7 The `end_me()` function

It will call `schedule` to transfer control to another thread, unless it is last thread. If it is the last thread, it will clean up all memory allocations, and `exit`.

Chapter 2

Data Structure Index

2.1 Data Structures

Here are the data structures with brief descriptions:

PriorityQueueEntryS	
An Entry in the Priority Queue	9
PriorityQueueS	
An Priority Queue for ProgramControlBlocks	10
ProgramControlBlockS	
Holds information about each thread in the system	11
SemaphoreS	
A Semaphore	11
thread_struct	12

Chapter 3

File Index

3.1 File List

Here is a list of all documented files with brief descriptions:

interface.h	Scheduler Interface definitions	13
scheduler.h	Scheduler function and data structures	14

Chapter 4


Data Structure Documentation

4.1 PriorityQueueEntryS Struct Reference

An Entry in the Priority Queue.

```
#include <scheduler.h>
```

Collaboration diagram for PriorityQueueEntryS:



Data Fields

- int **priority**
priority of the current entry
- struct **PriorityQueueEntryS** * **next**
pointer to data
- **ProgramControlBlock** * **block**
pointer to next entry, or Null if this is the last entry

4.1.1 Detailed Description

An Entry in the Priority Queue.

The documentation for this struct was generated from the following file:


- [scheduler.h](#)

4.2 PriorityQueueS Struct Reference

An Priority Queue for ProgramControlBlocks.

```
#include <scheduler.h>
```

Collaboration diagram for PriorityQueueS:



Data Fields

- [PriorityQueueEntry](#) * **head**
head of the queue
- [PriorityQueueEntry](#) * **tail**
tail of the queue

4.2.1 Detailed Description

An Priority Queue for ProgramControlBlocks.

The documentation for this struct was generated from the following file:

- [scheduler.h](#)

4.3 ProgramControlBlockS Struct Reference

Holds information about each thread in the system.

```
#include <scheduler.h>
```

Data Fields

- int [thread_id](#)
Unique id of the thread.
- int [cpu_burst_length](#)
How long has the current cpu burst lasted.
- bool [started](#)
Has the thread been started (specifically has it called startup_check() yet)
- int [arrival_time](#)
The [wall_clock](#) time whe the thread requested the current activity.
- int [remaining_time](#)
How much longer in ticks do the current activity require.
- int [remaining_time_quanta](#)
How many ticks remain in the threads CPU time quanta allocation.
- int [priority_level](#)
The priority level of the thread (only used in MLFQ)
- pthread_cond_t [allow_to_run](#)
A conditional used to block and unblock execution of this thread.

4.3.1 Detailed Description

Holds information about each thread in the system.

The documentation for this struct was generated from the following file:


- [scheduler.h](#)

4.4 SemaphoreS Struct Reference

A Semaphore.

```
#include <scheduler.h>
```

Collaboration diagram for SemaphoreS:



Data Fields

- `int value`
the current value of the semaphore
- `PriorityQueue blocked`
a queue for holding threads blocks by this semaphore

4.4.1 Detailed Description

A Semaphore.

The documentation for this struct was generated from the following file:

- `scheduler.h`

4.5 thread_struct Struct Reference

Data Fields

- `pthread_t p_t`
- `int tid`
- `char line [MAX_LINE_LEN]`

The documentation for this struct was generated from the following file:

- `main.c`


Chapter 5

File Documentation

5.1 interface.h File Reference

Scheduler Interface definitions.

This graph shows which files directly or indirectly include this file:



Macros

- `#define MAX_NUM_SEM 10`

Enumerations

- `enum sch_type { SCH_FCFS = 0 , SCH_SRTF = 1 , SCH_MLFQ = 2 }`

Functions

- `void init_scheduler (enum sch_type scheduler_type, int thread_count)`
- `int cpu_me (float current_time, int tid, int remaining_time)`
- `int io_me (float current_time, int tid, int duration)`
- `int P (float current_time, int tid, int sem_id)`
- `int V (float current_time, int tid, int sem_id)`
- `void end_me (int tid)`


5.1.1 Detailed Description

Scheduler Interface definitions.

5.2 scheduler.h File Reference

Scheduler function and data structures.

```
#include <stdio.h>
#include <stdlib.h>
#include <stdbool.h>
#include <math.h>
#include <limits.h>
#include <pthread.h>
#include <semaphore.h>
#include <assert.h>
#include "interface.h"
Include dependency graph for scheduler.h:
```



Data Structures

- struct [ProgramControlBlockS](#)
Holds information about each thread in the system.
- struct [PriorityQueueEntryS](#)
An Entry in the Priority Queue.
- struct [PriorityQueueS](#)
An Priority Queue for ProgramControlBlocks.
- struct [SemaphoreS](#)
A Semaphore.

Macros

- #define [MLFQ_LEVELS](#) sizeof(MLFQ_TIME_QUANTUM) / sizeof(int)
The number of levels to use in MLFQ.

Typedefs

- typedef struct [ProgramControlBlockS](#) [ProgramControlBlock](#)
Shorthand for struct [ProgramControlBlockS](#).
- typedef struct [PriorityQueueEntryS](#) [PriorityQueueEntry](#)
Shorthand for struct [PriorityQueueEntryS](#).
- typedef struct [PriorityQueueS](#) [PriorityQueue](#)
Shorthand for struct [PriorityQueueS](#).
- typedef struct [SemaphoreS](#) [Semaphore](#)
Shorthand for struct [SemaphoreS](#).

Functions

- void `PriorityQueue_init` (`PriorityQueue *queue`)
Initialize the PriorityQueue.
- void `PriorityQueue_enqueue` (`PriorityQueue *queue`, `ProgramControlBlock *block`, `int priority`)
Enqueue a ProgramControlBlock onto the queue according to priority.
- `ProgramControlBlock * PriorityQueue_dequeue` (`PriorityQueue *queue`)
Returns the first entry from the queue and removes it from the queue.
- `ProgramControlBlock * PriorityQueue_peek` (`PriorityQueue *queue`)
Returns the first entry from the queue without removing it from the queue.
- bool `PriorityQueue_is_empty` (`PriorityQueue *queue`)
Returns whether the queue is empty or not.
- void `PriorityQueue_destroy` (`PriorityQueue *queue`)
Removes and deallocate all entries in the queue.
- void `request_P` (`int arrival_time`, `int thread_id`, `int sem_id`)
Semaphore post implementation.
- void `request_V` (`int arrival_time`, `int thread_id`, `int sem_id`)
Semaphore wait implementation.
- void `request_cpu` (`int arrival_time`, `int thread_id`, `int remaining_time`)
Request to use the CPU.
- void `request_io` (`int arrival_time`, `int thread_id`, `int duration`)
Request to use IO.
- void `start_check` (`int arrival_time`, `int thread_id`)
Performs startup synchronization.
- void `schedule` (`int thread_id`)
Performs scheduling operations.
- void `tick` (`int num_ticks`)
Ticks the global clock, this should update current IO operations as well.
- void `add_to_ready_queue` (`ProgramControlBlock *block`)
Add the specified ProcessControlBlock to the ready queue.
- `ProgramControlBlock * pick_from_ready_queue` ()
Gets the next process in line to be run in the ready queue.
- `ProgramControlBlock * pick_from_arrival_queue` ()
Gets the next process to arrive from the arrival queue.
- `ProgramControlBlock * pick_from_io_queue` ()
Gets the a process that has just completed its IO operation. If such a process is available.
- int `next_arrival` ()
Gets how many ticks until the next process arrives.
- int `next_io_queue_completion` ()
Gets how many ticks until the current IO operation is finished.
- void `print_ready_queue` ()
DEBUG: print the ready queue.
- void `print_future_queue` ()
DEBUG: print the future queue.
- void `print_io_queue` ()
DEBUG: print the io queue.
- void `clean_up` ()
Cleans up before exit.

Variables

- int `num_processes`
total number of processes
- int `num_started_processes`
number of started processes
- int `num_running_processes`
number of running processes
- `ProgramControlBlock` * `program_control_blocks`
Array holding ProgramControlBlocks of each thread.
- int `wall_clock`
global simulation time
- `pthread_mutex_t` `execution_lock`
Mutex to ensure only one thread is active at any time.
- `PriorityQueue` `arrival_queue`
the arrival queue
- `PriorityQueue` * `ready_queue`
the Ready queue(s)
- `PriorityQueue` `io_queue`
the IO Queue
- `Semaphore` * `semaphores`
all Semaphores
- enum `sch_type` `scheduler_type`
The scheduler type.
- `ProgramControlBlock` * `current_cpu_program`
The currently executing thread on the cpu.

5.2.1 Detailed Description

Scheduler function and data structures.

Author

Jinyu Liu (jz16359@psu.edu)

Hongyu Guo (hbg5147@psu.edu)

Date

2022-10-10

5.2.2 Macro Definition Documentation

5.2.2.1 MLFQ_LEVELS

```
#define MLFQ_LEVELS sizeof(MLFQ_TIME_QUANTUM) / sizeof(int)
```

The number of levels to use in MLFQ.

Calculated from the length of the array `MLFQ_TIME_QUANTUM`

5.2.3 Function Documentation

5.2.3.1 add_to_ready_queue()

```
void add_to_ready_queue (
    ProgramControlBlock * block )
```


Add the specified ProcessControlBlock to the ready queue.

Exact implementation is [scheduler_type](#) dependent.

Parameters

<i>block</i>	the block to enqueue into the ready queue
--------------	---

Here is the caller graph for this function:



5.2.3.2 clean_up()

```
void clean_up ( )
```

Cleans up before exit.

Deallocates any memory allocated and deinitialize all data structures. This should be called by the last thread just before exiting `end_me()` Here is the call graph for this function:



5.2.3.3 next_arrival()


```
int next_arrival ( )
```

Gets how many ticks until the next process arrives.


Returns

int ticks until the next process arrives, -1 if no process is available

Here is the call graph for this function:



Here is the caller graph for this function:



5.2.3.4 next_io_queue_completion()


```
int next_io_queue_completion ( )
```

Gets how many ticks until the current IO operation is finished.


Returns

int ticks until the current IO operation finishes, -1 if no IO is in progress

Here is the call graph for this function:



Here is the caller graph for this function:

**5.2.3.5 pick_from_arrival_queue()**


```
ProgramControlBlock* pick_from_arrival_queue ( )
```

Gets the next process to arrive from the arrival queue.


Returns

ProgramControlBlock* the next process to arrive from the arrival queue.

Here is the call graph for this function:



Here is the caller graph for this function:



5.2.3.6 `pick_from_io_queue()`


```
ProgramControlBlock* pick_from_io_queue ( )
```

Gets the a process that has just completed its IO operation. If such a process is available.


Returns

ProgramControlBlock* a process that just completed IO, if available. NULL otherwise

Here is the call graph for this function:



Here is the caller graph for this function:



5.2.3.7 pick_from_ready_queue()

```
ProgramControlBlock* pick_from_ready_queue ( )
```


Gets the next process in line the be run in the ready queue.

The behavior of this function is dependent on the [scheduler_type](#).

Returns

ProgramControlBlock*

Here is the caller graph for this function:



5.2.3.8 PriorityQueue_dequeue()

```
ProgramControlBlock* PriorityQueue_dequeue (
    PriorityQueue * queue )
```

Returns the first entry from the queue and removes it from the queue.

Do not call this function with an empty queue. Use [PriorityQueue_is_empty\(\)](#) first to make sure queue is not empty.

Parameters

<i>queue</i>	the queue to operate on
--------------	-------------------------


Returns

ProgramControlBlock* the first entry in the queue

Here is the call graph for this function:



Here is the caller graph for this function:

**5.2.3.9 PriorityQueue_destroy()**

```
void PriorityQueue_destroy (
    PriorityQueue * queue )
```


Removes and deallocate all entries in the queue.

Note this does not deallocate the blocks, just the PriorityQueueEntries


Parameters

<i>queue</i>	the queue to destroy
--------------	----------------------

Here is the call graph for this function:



Here is the caller graph for this function:



5.2.3.10 PriorityQueue_enqueue()

```

void PriorityQueue_enqueue (
    PriorityQueue * queue,
    ProgramControlBlock * block,
    int priority )
  
```


Enqueue a ProgramControlBlock onto the queue according to priority.

Ties are broken using the block's `thread_id` with lower thread ids getting higher priority.

Parameters

<i>queue</i>	the queue to operate on
<i>block</i>	the block to enqueue
<i>priority</i>	the priority to assign to this block

Here is the caller graph for this function:



5.2.3.11 PriorityQueue_init()

```
void PriorityQueue_init (
    PriorityQueue * queue )
```

Initialize the PriorityQueue.

This should be called before any other operations are made to the queue

Parameters

<i>queue</i>	queue to operate on
--------------	---------------------

5.2.3.12 PriorityQueue_is_empty()

```
bool PriorityQueue_is_empty (
    PriorityQueue * queue )
```

Returns wether the queue is empty or not.

This should be called to ensure the queue is not empty before calling [PriorityQueue_dequeue\(\)](#) or [PriorityQueue_peek\(\)](#)

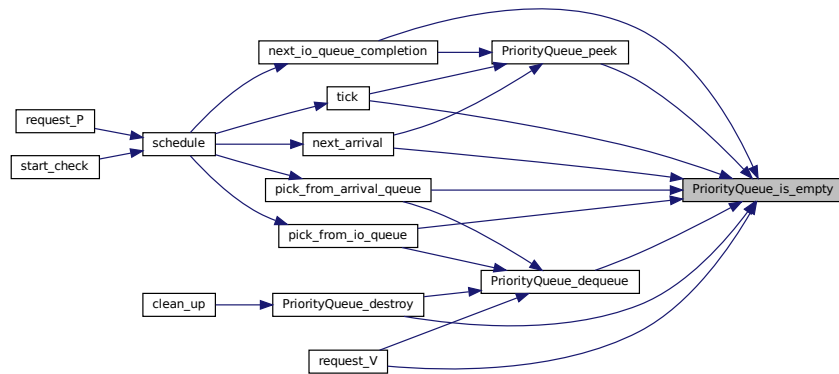
Parameters

<i>queue</i>	the queue to operate on
--------------	-------------------------

Returns

true the queue is empty
false the queue has at least one entry

Here is the caller graph for this function:

**5.2.3.13 PriorityQueue_peek()**

```
ProgramControlBlock* PriorityQueue_peek (
    PriorityQueue * queue )
```

Returns the first entry from the queue without removing it from the queue.

Do not call this function with an empty queue. Use [PriorityQueue_is_empty\(\)](#) first to make sure queue is not empty.

Parameters

<i>queue</i>	the queue to operate on
--------------	-------------------------


Returns

ProgramControlBlock* the first entry in the queue

Here is the call graph for this function:



Here is the caller graph for this function:



5.2.3.14 request_cpu()

```

void request_cpu (
    int arrival_time,
    int thread_id,
    int remaining_time )

```

Request to use the CPU.

Parameters

<i>arrival_time</i>	wall_clock time of arrival
<i>thread_id</i>	process requesting the CPU
<i>remaining_time</i>	how many more consecutive ticks will this process need the cpu

Here is the call graph for this function:



5.2.3.15 request_io()

```

void request_io (
    int arrival_time,
    int thread_id,
    int duration )

```

Request to use IO.

Parameters

<i>arrival_time</i>	wall_clock time of arrival
<i>thread_id</i>	process requesting the IO
<i>duration</i>	how long will the IO operation take.

Here is the call graph for this function:



5.2.3.16 request_P()

```

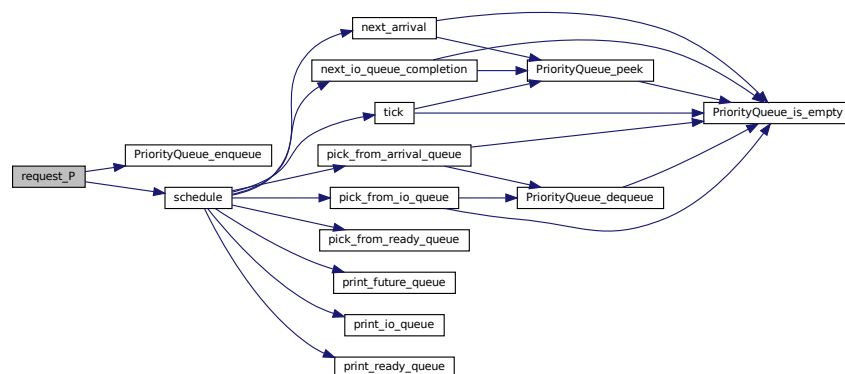
void request_P (
    int arrival_time,
    int thread_id,
    int sem_id )
  
```

Semaphore post implementation.

Parameters

<i>arrival_time</i>	The current wall_clock time
<i>thread_id</i>	The id of the calling thread
<i>sem_id</i>	The id of the semaphore to operate on

Here is the call graph for this function:



5.2.3.17 request_V()


```
void request_V (
    int arrival_time,
    int thread_id,
    int sem_id )
```

Semaphore wait implementation.

Parameters

<i>arrival_time</i>	The current wall_clock time
<i>thread_id</i>	The id of the calling thread
<i>sem_id</i>	The id of the semaphore to operate on

Here is the call graph for this function:



5.2.3.18 schedule()

```
void schedule (
    int thread_id )
```

Performs scheduling operations.

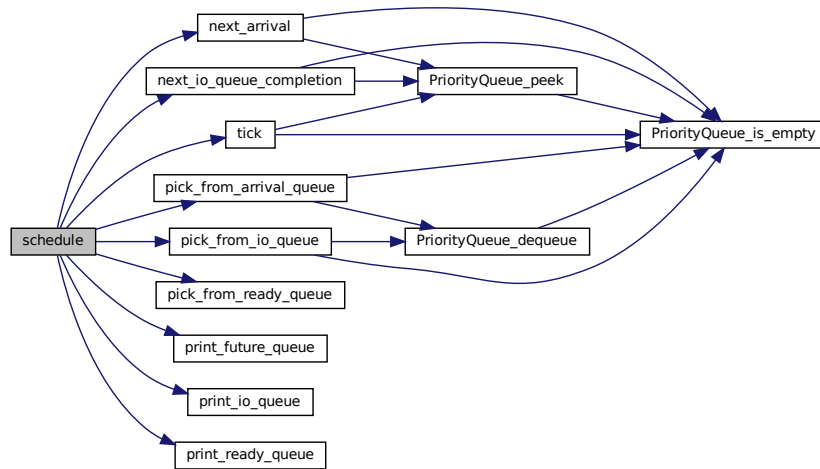
This program should be called in `cpu_me`, `io_me`, `V`, and `end_me` as long as there is still more operations to perform.

If it is decided that this thread should continue to execute then this function returns immediately, otherwise it will signal the thread to execute next and then block itself until another call to `schedule` decides that this thread should run next.


Parameters

<i>thread_id</i>	the id of the calling thread
------------------	------------------------------

Here is the call graph for this function:



Here is the caller graph for this function:



5.2.3.19 start_check()

```

void start_check (
    int arrival_time,
    int thread_id )

```

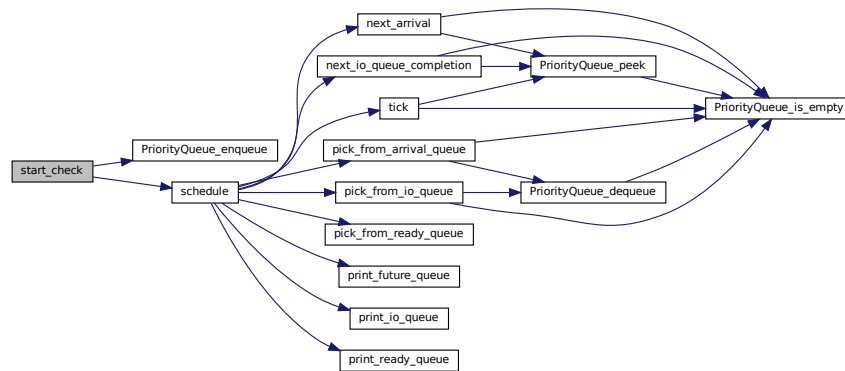
Performs startup synchronization.

Does nothing once the thread has been started. Call this before doing anything else in `cpu_me()`, `io_me()`, `P()`, `V()`.

Parameters

<i>arrival_time</i>	current wall_time
<i>thread_id</i>	id of the calling thread

Here is the call graph for this function:



5.2.3.20 tick()


```
void tick (
    int num_ticks )
```

Ticks the global clock, this should update current IO operations as well.


Parameters

<i>num_ticks</i>	number of ticks to skip forward
------------------	---------------------------------

Here is the call graph for this function:



Here is the caller graph for this function:



5.2.4 Variable Documentation

5.2.4.1 arrival_queue

```
PriorityQueue arrival_queue [extern]
```

the arrival queue

This queue holds all threads that will arrive in the future according to the current simulation time(wall_clock).

[pick_from_arrival_queue\(\)](#) and [next_arrival\(\)](#) operate on this queue.

5.2.4.2 current_cpu_program

```
ProgramControlBlock* current_cpu_program [extern]
```

The currently executing thread on the cpu.

This is null if cpu is idle.

5.2.4.3 execution_lock

```
pthread_mutex_t execution_lock [extern]
```

Mutex to ensure only one thread is active at any time.

All threads acquire this in their first call to [startup_check\(\)](#). They will temporarily relinquish it when blocked by their respective `ProgramControlBlockS::allowed_to_run` conditional. Finally they release it near the end of [end_me\(\)](#)

5.2.4.4 io_queue

```
PriorityQueue io_queue [extern]
```

the IO Queue

Holds threads that are performing or waiting to perform io. The head is the thread that is currently performing IO while all other threads are waiting

5.2.4.5 ready_queue

```
PriorityQueue* ready_queue [extern]
```

the Ready queue(s)

A set of queues to hold threads waiting for their turn on the cpu. The exact operations of this queue is dependent on the scheduler type

[add_to_ready_queue\(\)](#) and [pick_from_ready_queue\(\)](#) operates on this queue

5.2.4.6 scheduler_type

```
enum sch_type scheduler_type [extern]
```

The scheduler type.

as given in [init_scheduler\(\)](#)

5.2.4.7 wall_clock

```
int wall_clock [extern]
```

global simulation time

Do not change this value except using the [tick\(\)](#) function

Index

- add_to_ready_queue
 - [scheduler.h, 17](#)
- arrival_queue
 - [scheduler.h, 31](#)
- clean_up
 - [scheduler.h, 17](#)
- current_cpu_program
 - [scheduler.h, 31](#)
- execution_lock
 - [scheduler.h, 31](#)
- interface.h, [13](#)
- io_queue
 - [scheduler.h, 31](#)
- MLFQ_LEVELS
 - [scheduler.h, 16](#)
- next_arrival
 - [scheduler.h, 17](#)
- next_io_queue_completion
 - [scheduler.h, 18](#)
- pick_from_arrival_queue
 - [scheduler.h, 19](#)
- pick_from_io_queue
 - [scheduler.h, 20](#)
- pick_from_ready_queue
 - [scheduler.h, 20](#)
- PriorityQueue_dequeue
 - [scheduler.h, 21](#)
- PriorityQueue_destroy
 - [scheduler.h, 22](#)
- PriorityQueue_enqueue
 - [scheduler.h, 23](#)
- PriorityQueue_init
 - [scheduler.h, 24](#)
- PriorityQueue_is_empty
 - [scheduler.h, 24](#)
- PriorityQueue_peek
 - [scheduler.h, 25](#)
- PriorityQueueEntryS, [9](#)
- PriorityQueueS, [10](#)
- ProgramControlBlockS, [11](#)
- ready_queue
 - [scheduler.h, 32](#)
- request_cpu
 - [scheduler.h, 26](#)
- request_io
 - [scheduler.h, 26](#)
- request_P
 - [scheduler.h, 27](#)
- request_V
 - [scheduler.h, 28](#)
- schedule
 - [scheduler.h, 28](#)
- scheduler.h, [14](#)
 - [add_to_ready_queue, 17](#)
 - [arrival_queue, 31](#)
 - [clean_up, 17](#)
 - [current_cpu_program, 31](#)
 - [execution_lock, 31](#)
 - [io_queue, 31](#)
 - [MLFQ_LEVELS, 16](#)
 - [next_arrival, 17](#)
 - [next_io_queue_completion, 18](#)
 - [pick_from_arrival_queue, 19](#)
 - [pick_from_io_queue, 20](#)
 - [pick_from_ready_queue, 20](#)
 - [PriorityQueue_dequeue, 21](#)
 - [PriorityQueue_destroy, 22](#)
 - [PriorityQueue_enqueue, 23](#)
 - [PriorityQueue_init, 24](#)
 - [PriorityQueue_is_empty, 24](#)
 - [PriorityQueue_peek, 25](#)
 - [ready_queue, 32](#)
 - [request_cpu, 26](#)
 - [request_io, 26](#)
 - [request_P, 27](#)
 - [request_V, 28](#)
 - [schedule, 28](#)
 - [scheduler_type, 32](#)
 - [start_check, 29](#)
 - [tick, 30](#)
 - [wall_clock, 32](#)
- scheduler_type
 - [scheduler.h, 32](#)
- SemaphoreS, [11](#)
- start_check
 - [scheduler.h, 29](#)
- thread_struct, [12](#)
- tick
 - [scheduler.h, 30](#)
- wall_clock
 - [scheduler.h, 32](#)